Needs-based recommendation systems

Machine Learning Pipeline for Financial Needs Prediction and Product Recommendation

Authors: Peng Rao, Maxime d'Aviau de Ternay, Qiao Zhaochen, Shourya Vashistha, Arnau Nomen

Table of Contents

1. Project Overview
2. Exploratory Data Analysis
3. Data Preparation & Feature Engineering
4. ML Pipeline & Hyperparameter Optimization & Explainable AI (XAI)
5. Deep Learning Approach
6. Recommendations
7. Conclusion

Project Objectives

Dataset Structure

Data Quality Report

Target Variable Analysis

Feature Correlations with Targets

Key EDA Findings

Data Transformation Pipeline

Feature Transformation: Income & Wealth

Engineered Features (16 total)

ML Pipeline Overview

Baseline Models Evaluated

Baseline Results: AccumulationInvestment

Baseline Results: IncomeInvestment

Optuna Hyperparameter Tuning

Optimization Results

XAI Methods: SHAP, LIME & Permutation

Income Investment Model Insights

Accumulation Investment Model Insights

LIME: Local Interpretability Example

Deep Learning Model Architecture: DCN [1]

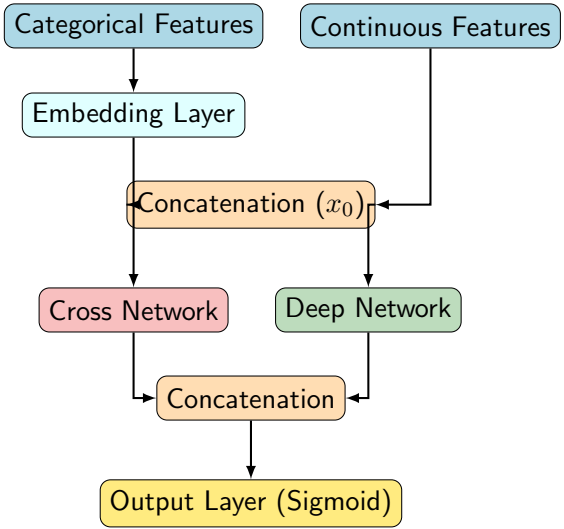

Multi-Task Learning Results

Task	Acc	Prec	Rec	F1	AUC
Income	0.783	0.852	0.544	0.664	0.767
Accum.	0.795	0.846	0.708	0.771	0.852

Architecture Highlights

- **Cross Network**: Applies explicit bounded-degree feature crossing efficiently.
- **Deep Network**: Learns implicit non-linear higher-order interactions.
- **Joint Training**: Both networks share the same concatenated input embeddings, optimising for dual investment targets.

Methodology: User-Item Matrix

We construct a user-item interaction matrix R using predicted probabilities from the supervised classifiers, penalised by risk misalignment.

Matrix Entry R_{ij}

$$R_{ij} = P_{\text{type}(j)}(i) \cdot \big(1 - \alpha \cdot |\text{RiskPropensity}(i) - \text{Risk}(j)|\big)$$

- $P_{\text{type}(j)}(i)$: Predicted probability of client i needing product j's type (Accumulation or Income).
- $\alpha \in [0, 1]$: Decay factor penalising risk misalignment.
  - **Economic Rationale**: Reflects the firm's risk aversion and regulatory suitability constraints (e.g., MiFID II).
  - **Trade-off**: Balances maximising expected fee revenue/conversion against minimising compliance risks and long-term customer churn from unsuitable investments.

Methodology: Models & Evaluation

We use model-based CF [2] leveraging matrix R:

- **Truncated SVD**: Factorises $R \approx U_k \Sigma_k V_k^T$ into low-dimensional latent factors to capture hidden preferences.
- **Autoencoder**: Learns non-linear representations by compressing and reconstructing the interaction matrix ($z_u = f_\theta(r_u)$, $\hat{r}_u = g_\phi(z_u)$).

Evaluation Metrics

- **RMSE**: Measures the model's reconstruction error.
- **Precision@k & Recall@k**: Evaluate the relevance of the top-k recommendations.
- **NDCG@k**: Assesses ranking quality, prioritising relevant items placed higher.

Experimental Results

Model	RMSE	Precision@3	Recall@3	NDCG@3
Truncated SVD	0.402	0.937	0.937	0.952
Autoencoder	**0.039**	**0.984**	**0.984**	**0.997**

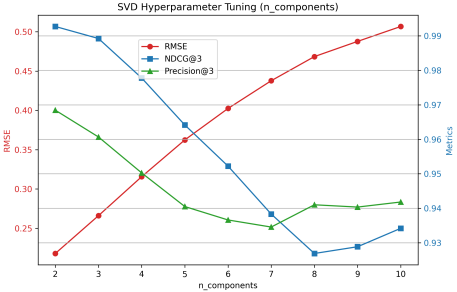

SVD Tuning

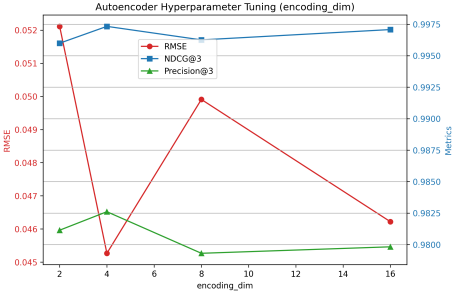

Autoencoder Tuning

Project Summary

References

[1] Weiyu Guo et al. "Explainable enterprise credit rating using deep feature crossing". en. In: *Expert Systems with Applications* 220 (June 2023), p. 119704.

[2] Marwa Sharaf et al. "A survey on recommendation systems for financial services". In: *Multimedia Tools and Applications* 81.12 (May 2022), pp. 16761–16781.